

Introduction to Parallel Computing

Byoung-Do (BD) Kim, PhD

Associate Chief Research Information Officer

Director, Center for Advanced Research Computing



USC

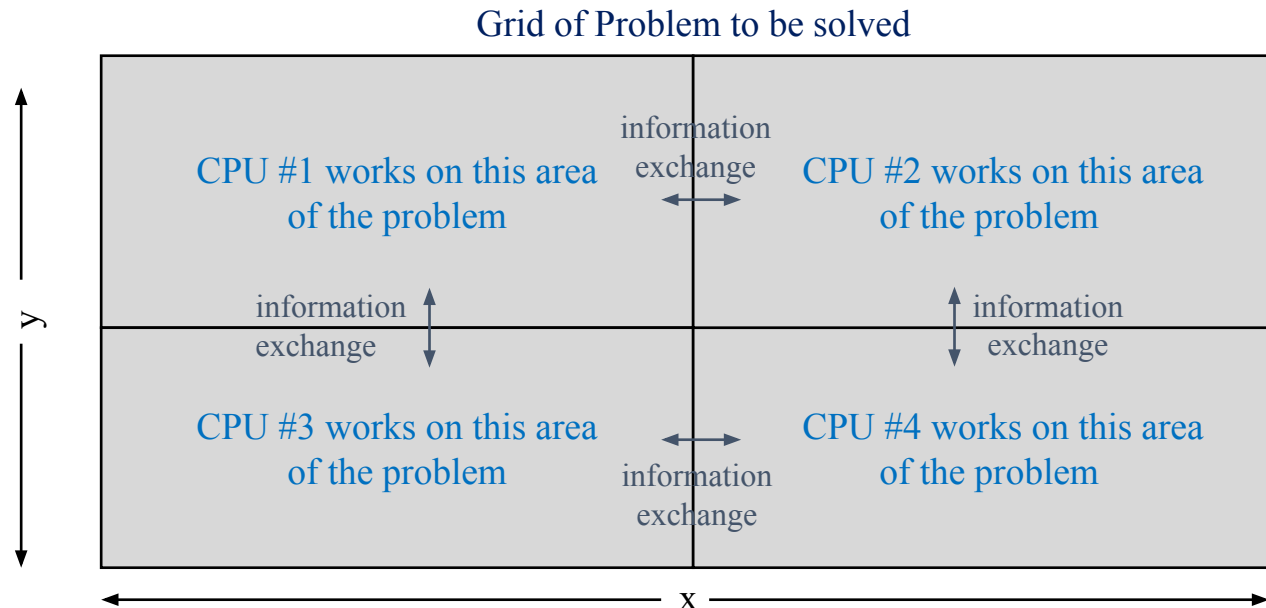
Advanced Research Computing
Enabling scientific breakthroughs at scale

Outline

1. Introduction
2. Theoretical background
3. Types of parallel computing systems
4. Programming models
5. Examples
6. Hands-on session w/examples

What is Parallel Computing?

- Parallel computing: use of multiple processors or computers working together on a common task.
 - Each processor works on its section of the problem
 - Processors can exchange information



Why Do Parallel Computing?

- Limits of single CPU computing
 - performance
 - available memory
- Parallel computing allows one to:
 - solve problems that don't fit on a single CPU
 - solve problems that can't be solved in a reasonable time
- We can solve...
 - larger problems
 - faster
 - more cases

Limits of Parallel Computing

- Theoretical Upper Limits
 - Amdahl's Law
- Practical Limits
 - Load balancing
 - Non-computational sections
 - Communication overhead
- Other Considerations
 - time to re-write code

Theoretical Upper Limit to Performance

- All parallel programs contain:
 - parallel sections (we hope!)
 - serial sections (unfortunately)
- Serial sections limit the parallel effectiveness
- Amdahl's Law states this formally

Amdahl's Law


- Amdahl's Law places a strict limit on the speedup that can be realized by using multiple processors.

- Speedup: $S = \frac{T_{serial}}{T_{parallel}}$

- Effect of multiple processors on speed up: $S = \frac{1}{f_s + \frac{f_p}{N}}$

where


- f_s = serial fraction of code
- f_p = parallel fraction of code
- N = number of processors



- Amdahl's law in multi-core era: <https://research.cs.wisc.edu/multifacet/amdahl/>

Illustration of Amdahl's Law


- It takes only a small fraction of serial content in a code to degrade the parallel performance.



From Wikipedia: Amdahl's Law
https://en.wikipedia.org/wiki/Amdahl%27s_law

Practical Limit: Amdahl's Law vs. Reality

- Amdahl's Law provides a theoretical upper limit on parallel speedup assuming that there are no costs for *communications*. In reality, communications will result in a further degradation of performance.




Practical Limit: Amdahl's Law vs. Reality

- In reality, the situation is even worse than predicted by Amdahl's Law due to:
 - Load balancing (waiting)
 - Scheduling (shared processors or memory)
 - Communications
 - I/O

Other Considerations


- In reality, the situation is even worse than predicted by Amdahl's Law
 - Scheduling (shared processors or memory)
 - Communications
 - I/O
- Writing effective parallel applications is difficult!
 - Load balance is important
 - Communication can limit parallel efficiency
 - Serial time can dominate
- Is it worth your time to rewrite your application?
 - Do the CPU requirements justify parallelization?
 - Will the code be used just once?

Shared vs. Distributed Memory



Shared memory:


- Single address space.
- All processors have access to a pool of shared memory.
- Methods of memory access :
 - Bus, Crossbar
- Programming model: OpenMP



Distributed memory:

- Each processor has its own local memory.
- Must do message passing to exchange data between processors.
- Methods of memory access :
 - Various topological interconnection
- Programming model: MPI

Multicore with Accelerators



- A limited number of processors N have access to a common pool of shared memory
- To use more than N processors requires data exchange over a network
- Communication details increasingly complex
 - Cache access
 - Main memory access
 - Quick Path / Hyper Transport socket connections
 - Node to node connection via network
- Load balancing critical for performance
- Requires specific libraries and compilers (CUDA, OpenCL, ACC, etc.)

Parallel Programming Models

- Data Parallelism
 - Each processor performs the same task on different data
- Task Parallelism
 - Each processor performs a different task on the same data (or on different data)
- Most applications fall between these two

Data Parallel Programming Example

- One code will run on 2 CPUs
- Program has array of data to be operated on 2 CPUs, array is split into two parts.

```
program:
...
if CPU=a then
    low_limit=1
    upper_limit=50
elseif CPU=b then
    low_limit=51
    upper_limit=100
end if
do I = low_limit,
upper_limit
    work on A(I)
end do
...
end program
```

CPU A

```
program:
...
low_limit=1
upper_limit=50
do I= low_limit,
upper_limit
    work on A(I)
end do
...
end program
```

CPU B

```
program:
...
low_limit=51
upper_limit=100
do I= low_limit,
upper_limit
    work on A(I)
end do
...
end program
```

Task Parallel Programming Example

- One code will run on 2 CPUs
- Program has 2 tasks (a and b) to be done by 2 CPUs

```
program.c:  
...  
initialize  
...  
if CPU=a then  
    do task a  
elseif CPU=b then  
    do task b  
end if  
...  
end program
```

CPU A

```
program.c:  
...  
initialize  
...  
do task a  
...  
end program
```

CPU B

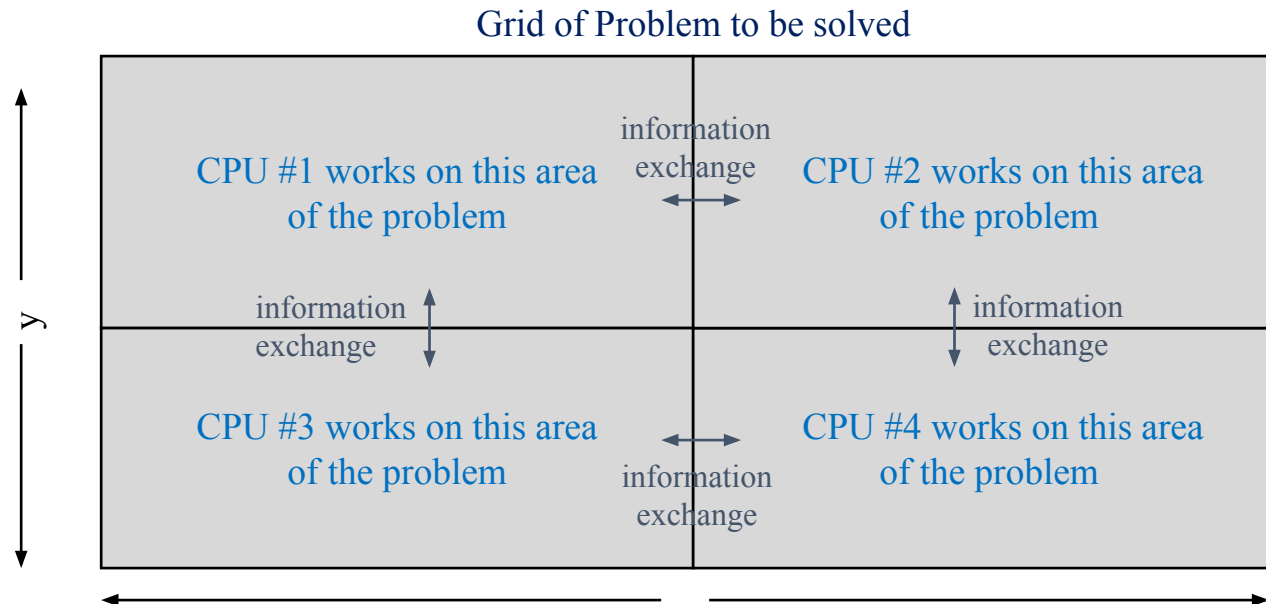
```
program.c:  
...  
initialize  
...  
do task b  
...  
end program
```


Single Program Multiple Data


- SPMD: dominant programming model for shared and distributed memory machines.
 - One source code is written
 - Code can have conditional execution based on which processor is executing the copy
 - All copies of code start simultaneously and communicate and sync with each other periodically
- MPMD: more general, and possible in hardware, but no system/programming software enables it

Data Decomposition

- For distributed memory systems, the ‘whole’ grid or sum of particles is decomposed to the individual processors
 - Each CPU works on its section of the problem
 - CPUs/Nodes can exchange information



SPMD Model



- Ideal programming model in multi-node system environment

Data Decomposition Example

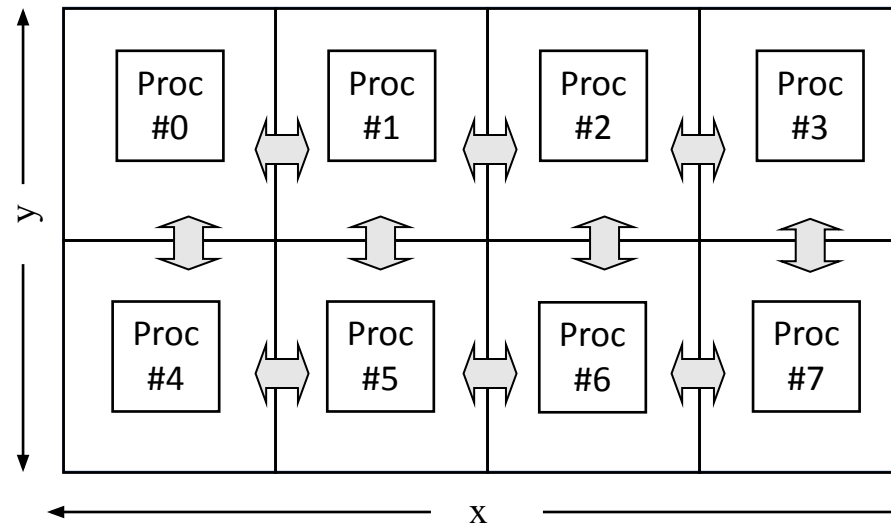
- Integrated 2-D propagation problem

Starting partial
differential equation:

$$\frac{\partial \Psi}{\partial t} = D \cdot \frac{\partial^2 \Psi}{\partial x^2} + B \cdot \frac{\partial^2 \Psi}{\partial y^2}$$

Finite Difference
Approximation:

$$\frac{f_{i,j}^{n+1} - f_{i,j}^n}{\Delta t} = D \cdot \frac{f_{i+1,j}^n - 2f_{i,j}^n + f_{i-1,j}^n}{\Delta x^2} + B \cdot \frac{f_{i,j+1}^n - 2f_{i,j}^n + f_{i,j-1}^n}{\Delta y^2}$$



MPI: Message Passing Interface

- Distributed memory programming
- Ideal for multi-node parallelization
- Can use with OpenMP for better scalability
- Distributed memory systems have separate address spaces for each processor
 - Local memory accessed faster than remote memory
 - **Data must be manually decomposed**
 - MPI is the standard for distributed memory programming

MPI Programming: Basics

Every MPI program needs these:


```
#include <mpi.h> /* the mpi include file */
int main(int argc, char *argv[])
{
    /* Initialize MPI */
    ierr = MPI_Init(&argc, &argv);
    /* How many total PEs are there */
    ierr = MPI_Comm_size(MPI_COMM_WORLD, &nPEs);
    /* What node am I (what is my rank? */
    ierr = MPI_Comm_rank(MPI_COMM_WORLD, &iam);
    ...
    ierr = MPI_Finalize();
}
```

MPI Example

```
#include
#include "mpi.h"

int main(int argc, char *argv[])
int argc;
char *argv[];
{
    int myid, numprocs;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    /* print out my rank and this run's PE size*/
    printf("Hello from %d\n",myid," of ",numprocs);
    MPI_Finalize();
}
```

Message Passing Communication



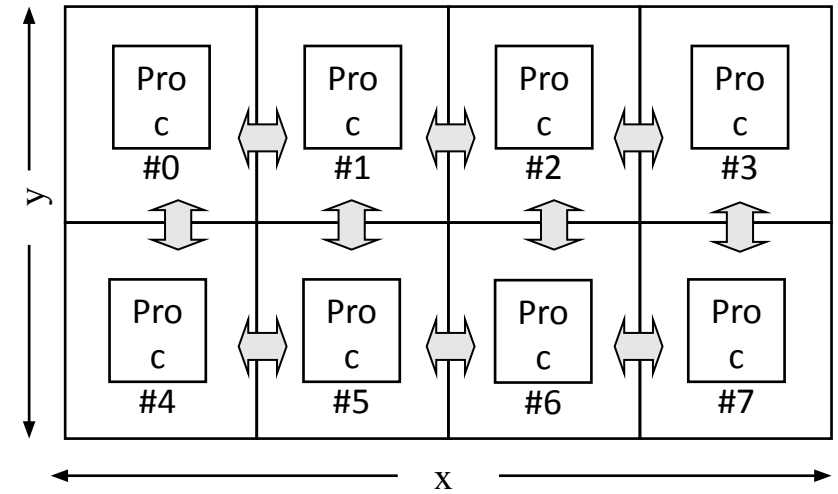
- Processes in message passing program communicate by passing messages
- Basic message passing primitives
 - `MPI_Send` (parameters list)
 - `MPI_Receive` (parameter list)
- These calls are blocking: the source processor issuing the send/receive cannot move to the next statement until the target processor issues the matching receive/send.

MPI Example: Send & Receive


```
#include "mpi.h"
/*****
This is a simple send/receive program in MPI
*****/
int main(int argc, char *argv[])
{
    int myid, numprocs, tag, source, destination, count, buffer ;
    MPI_Status status;
    MPI_Init(&argc, &argv) ;
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs) ;
    MPI_Comm_rank(MPI_COMM_WORLD, &myid) ;
    tag=1234;
    source=0;
    destination=1;
    count=1;
    if(myid == source){
        buffer=5678;
        MPI_Send(&buffer, count, MPI_INT, destination, tag, MPI_COMM_WORLD) ;
        printf("processor %d sent %d\n", myid, buffer) ;
    }
    if(myid == destination){
        MPI_Recv(&buffer, count, MPI_INT, source, tag, MPI_COMM_WORLD, &status) ;
        printf("processor %d got %d\n", myid, buffer) ;
    }
    MPI_Finalize() ;
}
```

Information Propagation between Processes


- How do we decide on what & how much of information should be passed along between decomposed computational domain?
 - based on PDE's *characteristics*



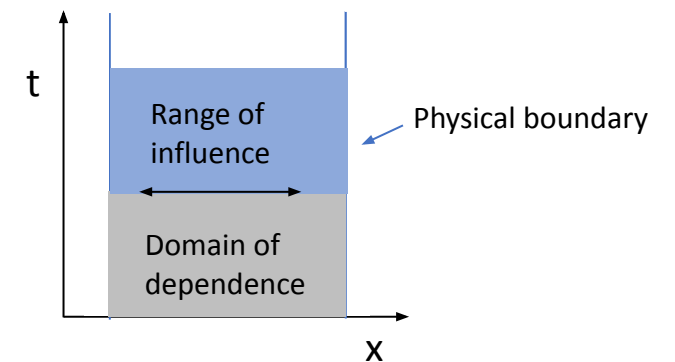
Elliptic PDE



Parabolic PDE



Hyperbolic PDE




* Information propagation speed:

Undefined

Infinite


Finite

Quick Overview of Parallel Computing



parallel
for ext

..
W



(time)

ter
ores

Hands-on session w/ examples

- Hello World
- Pi-calculation